

L16: Introduction to Rust (and More Parallel Programming)

Rachata Ausavarungnirun

(rachata.a@tggs.kmutnb.ac.th)

March 5th, 2020

Architecture Research Group

Software System Engineering

Thai-German Graduate School, KMUTNB

Systems and Parallel Programming

- Why is this typically hard?
- Systems programming
 - Hard to secure
 - Hard to multithread
- Parallel programs - Better Performance
 - Hard to detect data races
 - Hard to debug

Rust: Type Safety

- A well-defined program
 - No undefined behavior on all cases of execution
- A type safe language
 - Every program is well defined
- Is C type safe?
 - Then, why should we use C?

Python is not type safe, because of casting

(int \rightarrow float)

```
int main() {  
    int *a;  
    a = malloc(sizeof(int) * 2)  
    int *i;  
    i = a + 5  
    (*i)  
}
```

Performance Is King

- Performance differs across languages
 - Why? - Scala \rightarrow Assembly \rightarrow CPU (Different to Python)
- Computer executes an the assembly code
- Language somehow translate your program into assembly code
- And we have not even touch how hardware can be very different as well
 - See CPU vs. GPU
 - \hookrightarrow Something run well in CPU, but slow in GPU

Let's Compare Performance

- Assume a program that check if a string only consist of whitespace

Ruby Performance

- ```
class :: String
 def blank?
 /\A[[:space:]]*\z/ == self
 end
end
```
- You can run 964K of these per second  
(times)

# C Performance

- Let's Optimize this on a C code
  - I am going to borrow the code from [https://github.com/SamSaffron/fast\\_blank](https://github.com/SamSaffron/fast_blank)

*↳ Run every case in Parallel*

- You can run 10.5M iterations in 1 sec
- This is ~10 times faster

# Rust Performance

- **extern "C" fn fast\_blank(buf: Buf) -> bool {  
 buf.as\_slice().chars().all(|c| c.is\_whitespace())  
}**
- buf.as\_slice() gets the string slice
- .char() gets the iterator over each characters - *String → Character*
- Then the rest just check if there are whitespaces
- This is 11M iterations/sec



# Parallel Program in Rust

- Say I want to load many images from my input paths to all the images

(Function) (Variable Name) (Return)

- **fn** load\_images(paths: &[PathBuf]) -> Vec<Image> {  
    paths.iter()  
        .map(|path| {  
            Image::load(path)  
        })  
    .collect() }

- paths.iter() iterates over each path
- Then you load each path's image
- Create and return a vector of images

- **This is sequentially done** - Each path is independent

# Parallel Program in Rust

- I want to parallelize this
- **extern crate rayon** - Library to perform task in parallel  

```
fn load_images(paths: &[PathBuf]) -> Vec<Image> {
 paths.par_iter()
 .map(|path| {
 Image::load(path)
 })
 .collect() }
```
- rayon is a data-parallel library that convert sequential execution into parallel execution
  - <https://docs.rs/rayon/1.3.0/rayon/>
- paths.par\_iter() iterates over each path in parallel

# Detecting Data Races

- Rust compiler will tell you
- **fn** load\_images(paths: &[PathBuf]) -> Vec<Image> {  
(Val) **let mut** jpegs = 0;  
 paths.par\_iter()  
 .map(|path| {  
 **if** path.ends\_with("jpeg") { **jpegs += 1;** }  
 Image::load(path)  
 })  
 .collect();  
}
- This does not compile, why?

↳ Some Thread may read the old value, before adding it

# Detecting Data Races

- Rust compiler will tell you
- **fn** load\_images(paths: &[PathBuf]) -> Vec<Image> {  
    **let mut** jpegs = 0;  
    paths.par\_iter()  
        .map(|path| {  
            **if** path.ends\_with("jpeg") { jpegs += 1; }  
            Image::load(path)  
        })  
        .collect();  
}
- Note: let bind a value to a variable
- This will not compile, why?
  - Need to lock this (or use AtomicU32)

# Side Note on Atomic

- What is an atom?
- What is an atomic operation? - If something is wrong, go back to the beginning  
(Kill one of threads, and rewrite again)
- Why is this useful?

Ex

1000  $\rightarrow$  0  
ATM  $\rightarrow$  1000  
ATM2  $\rightarrow$  1000 } One of them get it

# Reading Rust Code

- fn defines a function

*Input variable name (mutable) and types* (64 bits unsigned integer)

- ```
fn gcd(mut n: u64, mut m: u64) -> u64 {  
    assert!(n != 0 && m != 0); This is a macro, ! asserts only if false  
    while m != 0 {  
        if m < n { Note no parenthesis for the condition check is ok  
            let t = m; m = n; n = t; Create a local variable. Type is inferred  
        }  
        m = m % n;  
    }  
    n You can also type return n, but the last line is the return value similar to Scala  
}
```

Generics

(Requirement)

Total Ordering (Whatever that can sort Ex. Integer, Object)

- `fn min<T: Ord>(a: T, b: T) -> T {
 if a <= b { a } else { b }
}`



Need to define how to
compare it

- Ord means that T, which is a generic type have total ordering


Enumeration and Sum Types

- `enum Option<T> {` *- Return Nothing or Some*
 `None,`
 `Some(T)`
 `}`
- `fn safe_div(n: i32, d: i32) -> Option<i32> {`
 `if d == 0 { return None; }`
 `return Some(n / d); }`
- `match safe_div(num, denom) {`
 `None => println!("No quotient."),`
 `Some(v) => println!("quotient is {}", v)`
 `}`

Catching Errors using Result<T,E>

- Result<T, E> is basically

- ```
enum Result<T, E> {
 Ok(T),
 Err(E),
}
```



- Let's use this to print the content of the directory and return the number of entries, or std::io::Error

# Catching Errors using Result<T,E>

- use std::path::Path;  
fn list\_directory(dir: &Path) -> std::io::Result<usize> {  
 let mut count = 0;  
 let entries = try!(std::fs::read\_dir(dir));  
 for entry\_or\_error in entries {  
 let entry = try!(entry\_or\_error);  
 println!("{:?}", entry.path());  
 count += 1;  
 }  
 return Ok(count);  
}
- fn main() {  
 let path = Path::new("/tmp");  
 let dir\_count = list\_directory(&path);  
 match dir\_count {  
 Ok(count) => println!("Total: {}", count),  
 Err(err) => println!("Error: {:?}", err) }  
}

# Memory Safety

- Rust provides three promises
- No null-pointer dereferencing
  - No Null value, use `Option<T>` or `Result<T, E>`
- No dangling pointers
  - What is a dangling pointer?
  - No garbage collectors as well
    - Use an ownership system
- No buffer overruns
  - No pointer arithmetic
  - Array in rust is not translated to pointers
  - Boundary checking at compile time

# In-Class Exercise 15

- Write a rust code that check if any number from n to m is prime
  - (optional) make this parallel